

Arabic Robot Research — Client Package

Egyptian Arabic voice robot — model selection bake-off + VPS research

Prepared 23 July 2026 · Formal reading guide (1 page)

ASR (default)

Whisper-Large-v3-Turbo-CT2

LLM (default)

Nile-Chat-4B

TTS (default)

VoiceTut-TTS

Host (default)

RunPod Secure · RTX 4090

Recommended reading order

- 1 Arabic_Robot_Research_Summary.pdf**
One-page overview of all decisions. Start here.
- 2 Combined_Families_Summary.pdf**
How ASR + LLM + TTS fit on one GPU (VRAM budget).
- 3 ASR / LLM / TTS selection reports**
Family bake-offs — recognition, brain, and voice.
- 4 VPS_Provider_Research_Report.pdf**
GPU cloud choice, capacity, and cost estimates.
- 5 TTS Voices/ (listen before freezing TTS)**
Confirm Egyptian dialect and code-switch quality.

Package contents

README.pdf (this page)
Arabic_Robot_Research_Summary.pdf
Combined_Families_Summary.pdf
ASR_Model_Selection_Report.pdf
LLM_Model_Selection_Report.pdf
TTS_Model_Selection_Report.pdf
VPS_Provider_Research_Report.pdf
TTS Voices/

- **VoiceTut-TTS.wav** ← recommended
- SILMA-TTS.wav (alternative)
- Chatterbox-Multilingual-V3.wav
- NMAAA-Egyptian-TTS.wav

Notes for sign-off

- TTS automated scores cover speed and resources only — please listen to the WAVs before freezing the voice.
- Quality LLM alternative: Qwen3-8B (int4). Accuracy ASR alternative: Whisper-Large-v3-CT2.
- Deploy on ≥ 24 GB VRAM (RTX 4090 class). Develop on RunPod Community; go live on RunPod Secure.
- GPU cloud prices change frequently — re-check provider dashboards before purchase.